Vysoká škola báňská - Technická univerzita Ostrava

Fakulta elektrotechniky a informatiky

Technologie databázových systémů 1
Semestrální projekt

Jméno: Phat Tran Dai

Osobní číslo: TRA0163   Datum: 21-05-2025

Obsah

Database Design

1.1. DD S01 L02

Explain the difference between the concept of data and information.

Data are raw facts or figures without context. For example, a number like 512, a string like 'hello world', or a timestamp like '2025-05-10 19:32:51' are all data. Information is data that has been processed or structured in a way that adds meaning. For instance, interpreting 512 as the duration (in seconds) of a video uploaded on '2025-05-10 19:32:51' transforms raw data into information.

1.2. DD S02 L02

Entities, instances, attributes and identifiers – describe in examples on your project.

- **Entity:** z_user (represents a user)

- **Instance:** A specific record in the table. Example: A user with user_id = 1234, username = 'jdoe', email = 'john.doe@example.com', and status = 'active'.

- **Attributes:** Properties of an entity. For z_user, attributes include username, email, first_name, last_name, profile_picture_id, etc.

- **Identifiers:** Unique attributes or their combinations that distinguish one instance from another. Example: user_id is the primary key of z_user, uniquely identifying each user.

1.3. DD S03 L01

- *Describe all relations in your database in English, including cardinality and membership obligation - each relation in two sentences.*

Name	Source	Cardinality Optionality	Target	Description
Category_Has_SubCategories	z_category	1..1 : 0..*	z_category	subcategory must refer to parent category, category can have many subcategories
Channel_Has_BannerImage	z_channel	0..* : 0..1	z_image_media	more channels can have the same banner image, channels can have a banner image
Channel_Has_Playlists	z_channel	0..1 : 0..*	z_playlist	playlist belongs to user, user can have many playlists
Channel_Has_ProfileImage	z_channel	0..* : 0..1	z_image_media	channel can have a profile image, image can be used as profile image by many channels
Channel_Has_Subscriptions	z_channel	1..1 : 0..*	z_subscription	subscriptions must refer to channel they target, channels can have many subscriptions
Channel_Has_Videos	z_channel	1..1 : 0..*	z_video	videos must refer to channel they belong to, channel can have many videos
Comment_Has_Reaction	z_comment	1..1 : 0..*	z_reaction	reaction must refer to comment it reacted to, comment can have many reactions
Comment_Has_SubComment	z_comment	1..1 : 0..*	z_comment	subcomment must refer to parent comment, comment can have many subcomments
Playlist_Has_Videos	z_playlist	1..1 : 0..*	z_playlist_video	playlist_video must refer to playlist it's in, playlist can have many playlist_videos
User_Has_Channels	z_user	1..1 : 0..*	z_channel	channel must refer to user it belongs to, user can have many channels
User_Has_Comments	z_user	1..1 : 0..*	z_comment	comment must refer to user it was written by, user can write many comments
User_Has_Playlists	z_user	1..1 : 0..*	z_playlist	user can have many playlists, playlist must refer to user
User_Has_ProfileImage	z_user	0..* : 0..1	z_image_media	image can be used as profile image by many users, user can have a profile image
User_Has_Reactions	z_user	1..1 : 0..*	z_reaction	reaction must refer to user, user can make make reactions
User_Has_Subscriptions	z_user	1..1 : 0..*	z_subscription	subscriptions must refer to a user, user can have many subscriptions
User_Has_VideoViews	z_user	1..1 : 0..*	z_video_view	video view must refer to user, user can make many video views
Video_Has_Comments	z_video	1..1 : 0..*	z_comment	comment must refer to exactly one video, video can have many comments
Video_Has_Reactions	z_video	1..1 : 0..*	z_reaction	reaction must refer to exactly one video, video can have many reactions
Video_Has_ThumbnailImage	z_video	0..* : 1..1	z_image_media	video must have a thumbnail, images can be thumbnails of many videos
Video_Has_VideoFile	z_video	0..1 : 1..1	z_video_media	video must have exactly one video file, video file can be the content of many videos
Video_Has_VideoViews	z_video	1..1 : 0..*	z_video_view	video view must refer to video, videos can have many views
Video_IsIn_Playlists	z_video	1..1 : 0..*	z_playlist_video	playlist_video must refer to the video, video can be many playlist_videos
z_video_category	z_video	0..* : 0..*	z_category	video can have many categories, categories can refer to many videos
User_OnChangeCreates_Audits	z_user	1..1 : 0..*	z_user_audit	user audit must refer to a user, user can have many audits
Comment_OnChangeCreates_Audits	z_comment	1..1 : 0..*	z_comment_audit	comment audit must refer to a comment, comment can have many audits
Video_OnChangeCreates_Audits	z_video	1..1 : 0..*	z_video_audit	video audits must refer to a video, video can have many audits

Table 1: Relations with cardinality and optionality

1.4. DD S03 L02

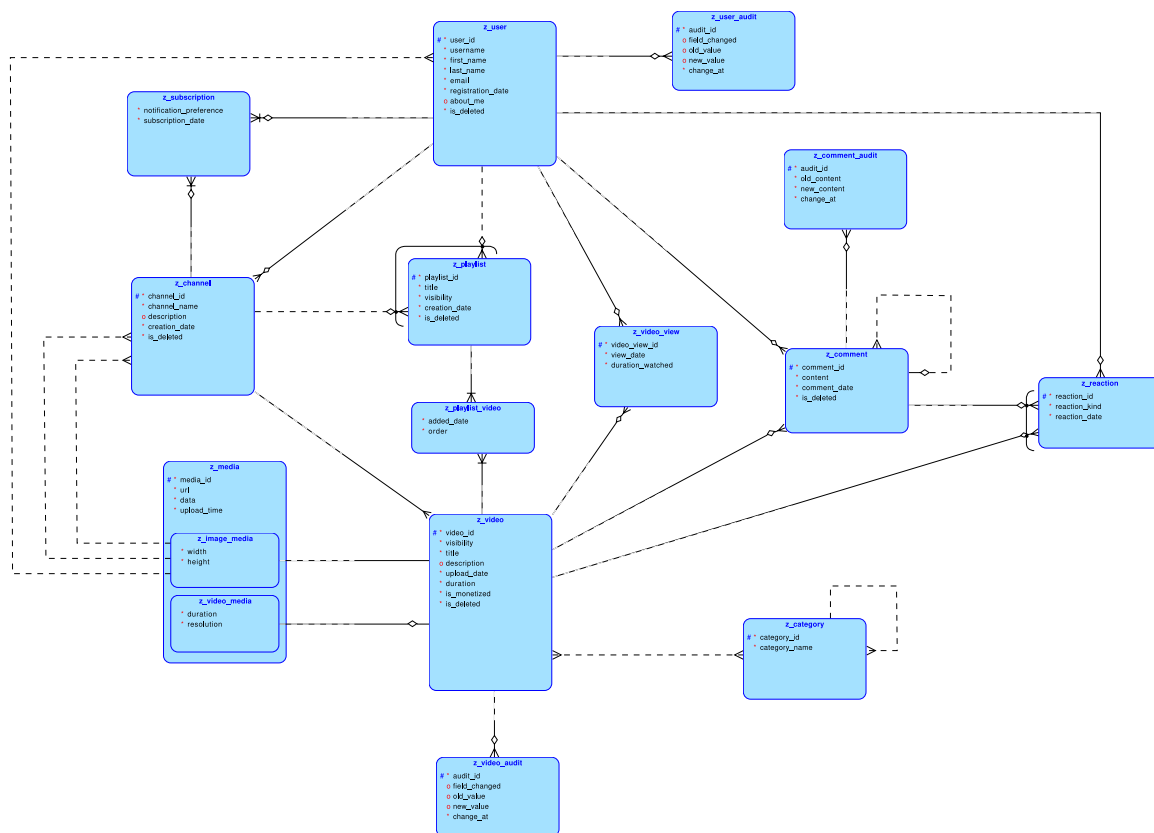*Draw an ER diagram according to conventions.*

Figure 1: ER Diagram

1.5. DD S30 L04

Matrix diagram with relationships, draw for your solution.

Should be read starting from the left.

	category	channel	comment	image_media	media	playlist	playlist_video	reaction	subscription	user	video	video_media	video_view
category	has subcategories										categorizes videos		
channel		—		has banner	references media	creates playlists			has subscribers	created by user	has videos		
comment			has replies					receives reactions		written by user	on video		
image_media		used by channel as pfp and banner		—	is subtype of					used in user for pfp			
media				is supertype of	—							is supertype of	
playlist		created by				—	contains videos			created by user			
playlist_video						belongs in	—				references video		
reaction			on comment					—		by user	on video		
subscription		to channel							—	by user			
user		creates channels	writes comments	has profile picture		creates playlists		creates reactions	makes subscriptions	—			makes views
video	is categorized	belongs to channel	has comments	has thumbnail image			is in playlists	receives reactions		viewed by users	—	has video file	has views
video_media					is subtype of						referenced by	—	
video_view										made by user	on video		—

Table 2: Relationship Matrix

1.6. DD S04 L01

Supertypes and subtypes – define at least one instance of a supertype and a subtype in your project.

The entity z_media is a generalization for all media. z_image_media and z_video_media are specializations that inherit media_id, data and created_at attributes from z_media but introduce additional attributes:

- z_image_media: width, height, format
- z_video_media: duration, resolution, codec

1.7. DD S04 L02

Description of business rules for your project.

- **Channel Ownership**
  - ‣ channel cannot be reassigned to another user post-creation
  - ‣ enforced by a BEFORE UPDATE trigger

- **Reaction Targeting**
  - ‣ reaction must target either a video or a comment, but never both
  - ‣ CHECK constraint

- **Playlist Creator Arc**
  - ‣ only one of user_id or channel_id can be set for a playlist
  - ‣ enforced by CHECK constraint

- **Soft Deletion**
  - ‣ deletions are logical via is_deleted, not physical

- **Non-Transferability**
  - ‣ certain foreign keys are immutable to preserve referential integrity (enforced with triggers)

- **Data Checks**
  - ‣ start time must start before end time, video cannot have negative views, etc.

1.8. DD S05 L01

Include at least one transferable and one non-transferable binding in your project.

- **Transferable Binding**
  - ‣ z_video.thumbnail_media_id → z_image_media.media_id is portable
  - ‣ supports reuse of shared images

- **Non-Transferable Binding**
  - ‣ z_comment.user_id is locked after creation via a BEFORE UPDATE trigger
  - ‣ non-transferable, immutable relationship

1.9. DD S05 L03

Have at least one M:N relationship without information and one M:N relationship with information in your project.

- **Without Information** - z_video_category links videos to categories. No extra attributes.
- **With Information** - z_playlist_video includes added_date and order, which add context beyond the M:N relation.

1.10. DD S06 L01

Incorporate at least one 1:N identifying relationship into your project, with the fact that the transferred foreign key will also be the key in the new table (identifying relationship).

- z_video_media.media_id is also the PK and a FK to z_media.media_id. This is a strong identifying relationship: each video media must be a media
- playlist_id and video_id of table z_playlist_video are identifying relationships, because the combination of them is the PK of z_playlist_video table

1.11. DD S06 L02-04

Have your schema in first normal form - no non-atomic attributes. In second normal form – no subkey bindings. In third normal form - no links between secondary attributes.

- *Normal Form Recap*
  - *1NF - no repeating groups or arrays; each attribute holds atomic (indivisible) values*
  - *2NF - must be in 1NF + no partial dependency on part of a composite key (i.e., every non-key attribute must depend on the whole primary key)*
  - *3NF - must be in 2NF + no transitive dependencies (i.e., non-key attributes must not depend on other non-key attributes)*

All the tables are in 3NF, because
- all non-key attributes functionally depend only on the primary key and
- there are no transitive dependecies.

1.12. DD S07 L01

Try to define ARC in your project (can be defined in ORACLE SQL Developer Data Modeler).

The table z_playlist uses an exclusive arc between user_id and channel_id. Only one of them may be non-null to indicate ownership. This is enforced via a CHECK constraint in DDL and supported in the modeler using arcs.

1.13. DD S07 L02

Try to define hierarchical and recursive relations in your project.

- z_category is hierarchical via parent_category_id (categories can nest).
- z_comment is recursive: replies point to another comment_id (enabling threaded discussions).

Both are actually hierarchical as well as recursive.

1.14. DD S07 L03

Describe how you record historical data in your system.

The system uses shadow/audit tables to preserve historical data. For example, every update to the z_comment.content or z_user is tracked in separate audit tables that store the previous and new value along with a timestamp and a reference to the original entity.

```
CREATE TABLE z_comment_audit (
  audit_id INTEGER PRIMARY KEY,
  comment_id INTEGER NOT NULL,
  old_content VARCHAR2(500),
  new_content VARCHAR2(500),
  change_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
  CONSTRAINT fk_comment_audit FOREIGN KEY (comment_id) REFERENCES z_comment(comment_id)
);

CREATE TABLE z_user_audit (
  audit_id INTEGER PRIMARY KEY,
  user_id INTEGER NOT NULL,
  field_changed VARCHAR2(50),
  old_value VARCHAR2(255),
  new_value VARCHAR2(255),
  change_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
  CONSTRAINT fk_user_audit FOREIGN KEY (user_id) REFERENCES z_user(user_id)
);
```

1.15. DD S09 L01

Demonstrate saving changes over time on your project.

When a tracked column such as z_user.email or z_comment.content is modified, the system logs the previous and new values (via triggers). This allows administrators or analysts to review the evolution of content or user states over time.

1.16. DD S09 L02

Try journaling in your project, i.e. saving past historical data (for example salary changes, workplace changes, etc.).

For user changes:

```sql
CREATE OR REPLACE TRIGGER trg_user_change
  BEFORE UPDATE OF status ON z_user
  FOR EACH ROW
BEGIN
  IF :OLD.email ≠ :NEW.email THEN
    INSERT INTO z_user_audit(user_id, field_changed, old_value, new_value, change_date)
    VALUES (:OLD.user_id, 'email', :OLD.email, :NEW.email);
  END IF;

  IF :OLD.first_name ≠ :NEW.first_name THEN
    INSERT INTO z_user_audit(user_id, field_changed, old_value, new_value, change_date)
    VALUES (:OLD.user_id, 'first_name', :OLD.first_name, :NEW.first_name);
  END IF;

  ... -- other attributes
END;
```

For comment content edits:

```sql
CREATE OR REPLACE TRIGGER trg_comment_content_change
  BEFORE UPDATE OF content ON z_comment
  FOR EACH ROW
BEGIN
  IF :OLD.content ≠ :NEW.content THEN
    INSERT INTO z_comment_audit(comment_id, old_content, new_content, change_date)
    VALUES (:OLD.comment_id, :OLD.content, :NEW.content, CURRENT_TIMESTAMP);
  END IF;
END;
```

1.17. DD S10 L01

Revise your design according to conventions for the readability of your schema.

ok.

1.18. DD S10 L02

Generic modeling – consider, possibly describe or use a generic model of data structures in your solution, how this approach is more advantageous compared to traditional data structure design methods.

Generic modeling refers to the design of data structures or database schemas in a way that is flexible, reusable, and abstracted from specific business cases. Rather than hardcoding entities or attributes for every use case (e.g., product_category, video_category, blog_category), a generic model creates unified, abstract structures that can serve many domains.

Example of domain specific non-general tables:

```sql
CREATE TABLE video_category (...);
CREATE TABLE blog_category (...);
```

In this design each use case has its own separate table.

Generic model would use something like this:

```sql
CREATE TABLE category (
  category_id INTEGER PRIMARY KEY,
  parent_category_id INTEGER REFERENCES category(category_id),
  name VARCHAR2(100)
);

CREATE TABLE entity_category (
  entity_type VARCHAR2(50),  -- e.g., 'video', 'blog'
  entity_id INTEGER,
  category_id INTEGER,
  PRIMARY KEY (entity_type, entity_id, category_id)
);
```

where entity_type may hold the name of the target table and the entity_id holds the id of the record it refers to.

This approach is advantageous when traditional design patterns like inheritance, polymorhphic arcs would nessecitate creations of large amounts of 'helper' tables. Generic modeling mitigates this problem of bloating the database model with such 'helper' tables by allowing storing field names and table names as field values of another table.

Although this pattern is somewhat of an anti-pattern it allows for scalable and simple models.

In the project there is z_user_audit table that holds historical data of the z_user table. In this table there is field_changed attribute which values are field names of the z_user table.

1.19. DD S11 L01

Describe examples of integrity constraints on your project for entities, bindings, attributes, and user-defined integrity.

- **Entity constraints**
  - ‣ z_user.user_id, z_video.video_id, etc. are primary keys → unique and not null
  - ‣ z_video.visibility has a CHECK constraint: must be 'draft', 'public', 'unlisted', or 'private'

- **Binding (relationship) constraints**
  - ‣ z_comment.video_id must exist in z_video (FOREIGN KEY)
  - ‣ z_subscription has a composite PK on (user_id, channel_id) to prevent duplicates

- **Attribute constraints**
  - ‣ z_user.email must be UNIQUE
  - ‣ z_video.duration must be a positive integer
  - ‣ z_reaction.reaction_kind must be 'like' or 'dislike'

- **User-defined constraints**
  - ‣ ARC logic:
    - – in z_reaction: either video_id or comment_id must be present, but not both
    - – in z_playlist, only one of user_id or channel_id can be non-null
    - – enforced via CHECK and TRIGGER

1.20. DD S11 L02-04

Generate a relational schema from your conceptual model and note the changes that have occurred in the schema and why.
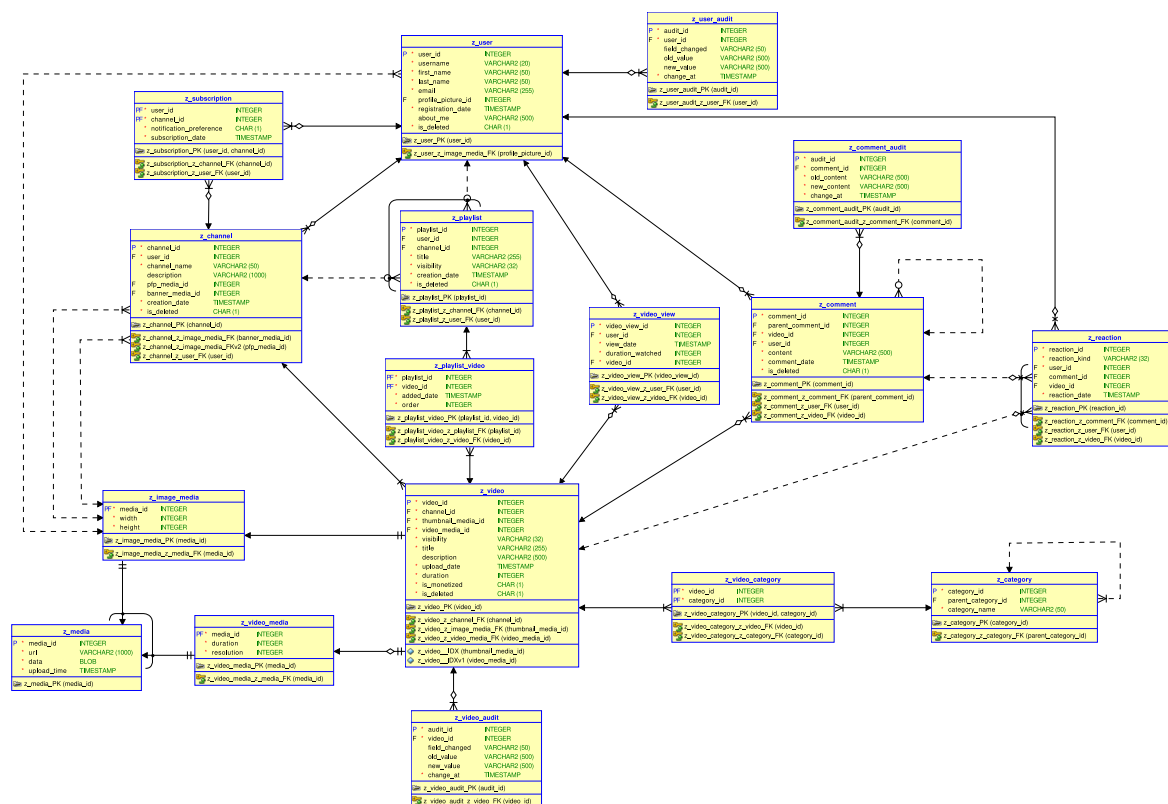

Figure 2: ER Diagram

Key changes from conceptual model

- Relations are modeled by adding foreign keys to tables.
- Many to many relations are modeled by adding an associative/junction table.
- Subtype implementation in SQL using shared key rather than classical inheritance
- ARC constraints require TRIGGERS due to SQL limitations on exclusive foreign key references

1.21. DD S15 L01

Write query for concatenate strings by pipes || , and CONCAT(). Write query with SELECT DISTINCT.

```sql
SELECT DISTINCT first_name || ' ' || last_name AS full_name
FROM z_user;

SELECT DISTINCT CONCAT(first_name, CONCAT(' ', last_name)) AS full_name
FROM z_user;
```

1.22. DD S16 L02

Write query with WHERE condition for selecting rows and use functions LOWER, UPPER, INITCAP.

```sql
SELECT * FROM z_user
WHERE LOWER(username) = 'phatt-23';

SELECT * FROM z_user
WHERE UPPER(status) = 'ACTIVE';

SELECT * FROM z_user
WHERE INITCAP(first_name) = 'John';
```

1.23. DD S16 L03

Write query with BETWEEN, AND, LIKE (%, _), IN(), IS NULL and IS NOT NULL.

```sql
SELECT * FROM z_video
WHERE duration BETWEEN 60 AND 300;

SELECT * FROM z_user
WHERE email LIKE '%.cz';

SELECT * FROM z_user
WHERE username LIKE '_ohn';

SELECT * FROM z_user
WHERE status IN ('active', 'suspended');

SELECT * FROM z_user
WHERE profile_picture_id IS NULL;

SELECT * FROM z_user
WHERE profile_picture_id IS NOT NULL;
```

1.24. DD S17 L01

Write query with AND, OR, NOT and precendence.

```sql
SELECT * FROM z_user
WHERE (status = 'active' AND is_deleted = '0')
   OR (status = 'suspended' AND NOT is_deleted = '1');
```

1.25. DD S17 L02

Write query with ORDER BY atr [ASC/DESC] - Sorting by using one or more attributes.

```sql
SELECT * FROM z_video
ORDER BY upload_date DESC;

SELECT * FROM z_video
ORDER BY visibility ASC, like_count DESC;
```

1.26. DD S17 L03

Write query with single row functions and column functions MIN, MAX, AVG, SUM, COUNT.

```sql
SELECT COUNT(*) AS total_users FROM z_user;

SELECT MIN(duration) AS shortest_video, MAX(duration) AS longest_video FROM z_video;

SELECT AVG(duration) AS average_duration FROM z_video;

SELECT SUM(view_count) AS total_views FROM z_video;
```

Programming with SQL

2.1. SQL S01 L01

- LOWER, UPPER, INITCAP
- CONCAT, SUBSTR, LENGTH, INSTR, LPAD, RPAD, TRIM, REPLACE
- Use virtual table DUAL

```sql
-- LOWER, UPPER, INITCAP
SELECT
    LOWER('HeLLo WorlD') AS lower,
    UPPER('HeLLo WorlD') AS upper,
    INITCAP('heLLo worLD') AS initcap
FROM dual;

-- CONCAT, SUBSTR, LENGTH, INSTR
SELECT
    CONCAT('Hello', 'World') AS concatenated,
    SUBSTR('abcdef', 2, 3) AS substr_example, -- 'bcd'
    LENGTH('Oracle') AS len,
    INSTR('banana', 'a') AS first_a_position
FROM dual;

-- LPAD, RPAD
SELECT
    LPAD('42', 5, '0') AS lpad_example,  -- '00042'
    RPAD('ok', 5, '.') AS rpad_example   -- 'ok...'
FROM dual;

-- TRIM, REPLACE
SELECT
    TRIM(' x ') AS trimmed,              -- 'x'
    REPLACE('abracadabra', 'a', 'X') AS replaced
FROM dual;
```

2.2. SQL S01 L02

ROUND, TRUNC round for two decimal places, whole thousands MOD

```sql
SELECT
    ROUND(123.4567, 2) AS round_2dec,    -- 123.46
    TRUNC(123.4567, 2) AS trunc_2dec     -- 123.45
FROM dual;

-- ROUND and TRUNC to whole thousands
SELECT
    ROUND(98765, -3) AS round_thousands,  -- 99000
    TRUNC(98765, -3) AS trunc_thousands   -- 98000
FROM dual;

-- MOD to get remainder
SELECT MOD(17, 5) AS remainder  -- 2
FROM dual;
```

2.3. SQL S01 L03

- MONTHS_BETWEEN, ADD_MONTHS, NEXT_DAY, LAST_DAY, ROUND, TRUNC
- System constant SYSDATE

```sql
-- MONTHS_BETWEEN, ADD_MONTHS, NEXT_DAY, LAST_DAY, ROUND, TRUNC
SELECT
    SYSDATE AS now,
    MONTHS_BETWEEN(DATE '2025-05-01', DATE '2025-03-01') AS months_between,
    ADD_MONTHS(SYSDATE, 2) AS two_months_later,
    NEXT_DAY(SYSDATE, 'MONDAY') AS next_monday,
    LAST_DAY(SYSDATE) AS last_day_of_month,
    ROUND(SYSDATE, 'MONTH') AS rounded_to_month,
    TRUNC(SYSDATE, 'MONTH') AS trunc_to_month
FROM dual;
```

2.4. SQL S02 L01

- TO_CHAR, TO_NUMBER, TO_DATE

```sql
-- TO_CHAR, TO_NUMBER, TO_DATE
SELECT
    TO_CHAR(SYSDATE, 'YYYY-MM-DD HH24:MI') AS formatted_date_varchar2,
    TO_NUMBER('12345.67', '99999.99') AS to_number_example,
    TO_DATE('2025-12-01', 'YYYY-MM-DD') AS to_date_example
FROM dual;
```

2.5. SQL S02 L02

- NVL, NVL2, NULLIF, COALESCE

```sql
-- NVL, NVL2, NULLIF, COALESCE
SELECT
    NVL(NULL, 'default') AS nvl_example,        -- 'default'
    NVL2(NULL, 'a', 'b') AS nvl2_null,          -- 'b'
    NVL2('X', 'a', 'b') AS nvl2_notnull,        -- 'a'
    NULLIF(100, 100) AS nullif_equal,           -- NULL
    COALESCE(NULL, NULL, 'first non-null') AS coalesce_example
FROM dual;
```

2.6. SQL S02 L03

- DECODE, CASE, IF-THEN-ELSE

```sql
-- DECODE, CASE, IF-THEN-ELSE equivalent (only CASE is SQL standard)
SELECT
    DECODE(2,            -- expr switch:
        1, 'one',        --   expr1 ⇒ 'one'
        2, 'two',        --   expr2 ⇒ 'two'
          'other'        --   _ ⇒ 'other'
    ) AS decode_example,
    CASE 3
        WHEN 1 THEN 'one'
        WHEN 2 THEN 'two'
        WHEN 4 THEN 'four'
        ELSE 'other' END
    AS case_example
FROM dual;
```

2.7. SQL S03 L01

- NATURAL JOIN, CROSS JOIN

```sql
SELECT *
FROM z_user
NATURAL JOIN z_playlist;   -- joins by common column names

SELECT *
FROM z_user
CROSS JOIN z_channel;  -- cartesian product, can't specify join clause
```

2.8. SQL S03 L02

- JOIN … USING(atr)
- JOIN … ON (joining condition)

```sql
SELECT *
FROM z_video_view
JOIN z_video USING(video_id);    -- USING(column1, column2, ...)

SELECT *
FROM z_video_view
JOIN z_video ON (z_video_view.video_id = z_video.video_id);
```

2.9. SQL S03 L03

- LEFT OUTER JOIN … ON ()
- RIGHT OUTER JOIN … ON ()
- FULL OUTER JOIN … ON ()

```sql
SELECT * FROM z_video v LEFT OUTER JOIN z_video_view vv USING(video_id);
-- same as
SELECT * FROM z_video v LEFT JOIN z_video_view vv USING(video_id);

SELECT * FROM z_video v RIGHT OUTER JOIN z_video_view vv USING(video_id);

SELECT * FROM z_video v FULL OUTER JOIN z_video_view vv USING(video_id);
-- same as
SELECT * FROM z_video v LEFT OUTER JOIN z_video_view vv USING(video_id)
UNION
SELECT * FROM z_video v RIGHT OUTER JOIN z_video_view vv USING(video_id);
```

2.10. SQL S03 L04

- Joining 2× of the same table with renaming (link between superiors and subordinates in one table)
- Hierarchical querying – tree structure of START WITH, CONNECT BY PRIOR, LEVEL

```sql
-- Self-join to link parent/child comments
SELECT
    a.comment_id AS parent_id,
    b.comment_id AS child_id
FROM z_comment a
JOIN z_comment b ON a.comment_id = b.parent_comment_id;

-- Hierarchical query for comment tree
SELECT
    LEVEL,
    comment_id,
    parent_comment_id,
    content
FROM z_comment
START WITH parent_comment_id IS NULL
CONNECT BY PRIOR comment_id = parent_comment_id;
```

```
DECLARE
    v_parent_comment_id INT := 12;
    v_user_id INT := 12;
    v_video_id INT := 1;
BEGIN
    FOR i IN 0..10 LOOP
        INSERT INTO z_comment (parent_comment_id, video_id, user_id, content)
        VALUES (
            (SELECT comment_id FROM z_comment WHERE comment_id = v_parent_comment_id),
            (SELECT video_id FROM z_video WHERE video_id = v_video_id),
            (SELECT user_id FROM z_user WHERE user_id = v_user_id),
            'hello, world! reply to comment ' || i
        );
    END LOOP;
END;


DECLARE
    v_level INT := 0;
    v_comment_id Z_COMMENT.comment_id%TYPE := 0;
    v_parent_comment_id Z_COMMENT.comment_id%TYPE := 0;
    v_content Z_COMMENT.content%TYPE := '';
BEGIN
    FOR v_rec IN (
        SELECT
            LEVEL,
            comment_id,
            parent_comment_id,
            content
        FROM z_comment
        START WITH parent_comment_id IS NULL
        CONNECT BY PRIOR comment_id = parent_comment_id
    ) LOOP
        FOR v_tab IN 1..v_rec.LEVEL LOOP
            DBMS_OUTPUT.PUT('    ');
        END LOOP;

        DBMS_OUTPUT.PUT_LINE(v_rec.comment_id || ' → ' || v_rec.content);
    END LOOP;
END;
```

2.11. SQL S04 L02

- AVG, COUNT, MIN, MAX, SUM, VARIANCE, STDDEV

```
BEGIN
    FOR rec IN (SELECT * FROM z_user CROSS JOIN z_video)
    LOOP
        INSERT INTO z_video_view(user_id, duration_watched, video_id)
        VALUES (rec.user_id, rec.duration / 2, rec.video_id);
    END LOOP;
END;

WITH t_video_view AS (
    SELECT
        video_id AS video_id,
        duration AS duration,
        COUNT(video_view_id) AS view_count
    FROM z_video JOIN z_video_view USING(video_id)
    GROUP BY video_id, duration
), t_reaction_count AS (
    SELECT
        (SELECT COUNT(*) FROM z_reaction r
         WHERE r.video_id = v.video_id
         AND reaction_kind = 'like') AS like_count,
        (SELECT COUNT(*) FROM z_reaction r
         WHERE r.video_id = v.video_id
         AND reaction_kind = 'dislike') AS dislike_count
    FROM z_video v
)
SELECT
    COUNT(*),
    AVG(view_count),
    MIN(duration),
    MAX(duration),
    SUM(like_count),
    AVG(dislike_count),
    VARIANCE(dislike_count),  -- Var(X) := E((X - mean) ** 2)
    SQRT(VARIANCE(dislike_count)) AS standard_deviation,  -- StdDev(X) := sqrt2(Var(X))
    STDDEV(dislike_count)
FROM t_video_view, t_reaction_count;
```

2.12. SQL S04 L03

- COUNT, COUNT(DISTINCT), NVL
- Difference between COUNT (*) a COUNT (attribute)
- Why using NVL for aggregation functions

```sql
SELECT
    COUNT(*),
    COUNT(thumbnail_media_id),
    COUNT(DISTINCT visibility)
FROM z_video;



WITH t AS (
    SELECT
        v.video_id AS video_id,
        NULLIF(COUNT(VIDEO_VIEW_ID), 0) AS view_count
    FROM z_video v
    LEFT JOIN z_video_view vv ON v.video_id = vv.video_id
    GROUP BY v.video_id
)
SELECT
    AVG(NVL(view_count, 0)),
    -- not the same AVG, null is disregarded by default
    AVG(view_count)
FROM t;
```

2.13. SQL S05 L01

- GROUP BY
- HAVING

```sql
SELECT
    channel_id,
    COUNT(*) AS total_videos
FROM z_video
GROUP BY channel_id;

SELECT
    channel_id,
    COUNT(*) AS total_videos
FROM z_video
GROUP BY channel_id
HAVING COUNT(*) > 1;
```

2.14. SQL S05 L02

- ROLLUP, CUBE, GROUPING SETS

```sql
-- GROUPING SETS
SELECT
    channel_id,
    visibility AS video_visibility,
    COUNT(*) AS video_count
FROM z_video
GROUP BY GROUPING SETS (
    (channel_id),
    (visibility)
);

-- equivalent to: (using UNION is much more inefficient)
SELECT
    channel_id,
    NULL AS video_visibility,
    COUNT(*) AS video_count
FROM z_video
GROUP BY channel_id

UNION ALL

SELECT
    NULL AS channel_id,
    visibility AS video_visibility,
    COUNT(*) AS video_count
FROM z_video
GROUP BY visibility;
```

```sql
-- GROUP BY CUBE
-- every combination, resulting in 2 ** n grouping sets
-- where n is number of columns in the cube list
SELECT
    channel_id,
    visibility AS video_visibility,
    COUNT(*) AS video_count
FROM z_video
GROUP BY CUBE(channel_id, visibility)
ORDER BY channel_id, video_visibility, video_count NULLS LAST;

-- is equivalent to this:
SELECT
    channel_id,
    visibility AS video_visibility,
    COUNT(*) AS video_count
FROM z_video
GROUP BY GROUPING SETS(
    (),
    (channel_id),
    (visibility),
    (channel_id, visibility)
)
ORDER BY channel_id, video_visibility, video_count NULLS LAST;
```

```sql
-- is equivalent to this:
SELECT *
FROM (
    SELECT NULL AS channel_id,
           NULL AS video_visibility,
           COUNT(*) AS video_count
    FROM z_video

    UNION ALL

    SELECT channel_id AS channel_id,
           NULL AS video_visibility,
           COUNT(*) AS video_count
    FROM z_video
    GROUP BY channel_id

    UNION ALL

    SELECT NULL AS channel_id,
           visibility AS video_visibility,
           COUNT(*) AS video_count
    FROM z_video
    GROUP BY visibility

    UNION ALL

    SELECT channel_id AS channel_id,
           visibility AS video_visibility,
           COUNT(*) AS video_count
    FROM z_video
    GROUP BY channel_id, visibility
)
ORDER BY channel_id, video_visibility, video_count NULLS FIRST;
```

```sql
-- GROUP BY ROLLUP
-- group by channel_id, visibility
-- with hierarchical subtotals per group and grand total being the sum of all subtotals
SELECT
    channel_id,
    visibility AS video_visibility,
    COUNT(*) AS video_count
FROM z_video
GROUP BY ROLLUP(channel_id, visibility)
ORDER BY channel_id, video_visibility, video_count NULLS FIRST;

-- equivalent to this:

SELECT
    channel_id,
    visibility AS video_visibility,
    COUNT(*) AS video_count
FROM z_video
GROUP BY GROUPING SETS(
    (channel_id, visibility),
    (channel_id),
    ()
)
ORDER BY channel_id, video_visibility, video_count NULLS FIRST;
```

```
-- equivalent to this:

SELECT * FROM (
    SELECT
        channel_id AS channel_id,
        visibility AS video_visibility,
        COUNT(*) AS video_count
    FROM z_video
    GROUP BY channel_id, visibility

    UNION ALL

    SELECT
        channel_id AS channel_id,
        NULL AS video_visibility,
        COUNT(*) AS video_count
    FROM z_video
    GROUP BY channel_id

    UNION ALL

    SELECT
        NULL AS channel_id,
        NULL AS video_visibility,
        COUNT(*) AS video_count
    FROM z_video
)
ORDER BY channel_id, video_visibility, video_count NULLS FIRST;
```

2.15. SQL S05 L03

- Multiple operatíons in SQL – UNION, UNION ALL, INTERSECT, MINUS
- ORDER BY for set operations

```sql
SELECT * FROM z_user;
SELECT * FROM z_user WHERE email LIKE '%yahoo.com';

SELECT username FROM z_user WHERE is_deleted = 0
UNION
SELECT channel_name FROM z_channel WHERE is_deleted = 0;

SELECT * FROM z_user WHERE email LIKE '%gmail.com'; -- 3
SELECT * FROM z_user WHERE first_name = 'Shelly'; -- 1

SELECT username FROM z_user WHERE first_name = 'Shelly'
INTERSECT -- 1
SELECT username FROM z_user WHERE email LIKE '%gmail.com';

SELECT username FROM z_user WHERE email LIKE '%gmail.com'
MINUS -- 2
SELECT username FROM z_user WHERE first_name = 'Shelly';

SELECT username FROM z_user
UNION ALL
SELECT channel_name FROM z_channel
ORDER BY 1; -- first item in the select list
-- ORDER BY username; -- could also be this though
```

2.16. SQL S06 L01

- Nested queries
  - ‣ Result as a single value
  - ‣ Multi-column subquery
  - ‣ EXISTS, NOT EXISTS

```sql
SELECT username
FROM z_user
WHERE user_id = (SELECT user_id FROM z_channel WHERE channel_name = 'Hicks Ltd');

SELECT
    v.channel_id,
    v.title
FROM z_video v
WHERE (v.channel_id, v.visibility) IN (
    SELECT p.channel_id, p.visibility FROM z_playlist p
);


-- users with at least one video
SELECT *
FROM z_user u
WHERE EXISTS (
    SELECT 1 FROM z_video v
    WHERE v.channel_id IN (
        SELECT c.channel_id
        FROM z_channel c
        WHERE c.user_id = u.user_id
    )
);
```

2.17. SQL S06 L02

- One-line subqueries

```sql
SELECT
    channel_name,
    (
        SELECT COUNT(*)
        FROM z_video v
        WHERE v.channel_id = c.channel_id
    ) AS video_count
FROM z_channel c;
```

2.18. SQL S06 L03

- Multi-line subqueries IN, ANY, ALL
- NULL values in subqueries

```sql
-- users with private playlists
SELECT username
FROM z_user
WHERE user_id IN (SELECT user_id FROM z_playlist WHERE visibility = 'private');


-- playlists that belong to deleted channels
SELECT playlist_id, channel_id
FROM z_playlist
WHERE channel_id = ANY (
    SELECT channel_id
    FROM z_channel
    WHERE is_deleted = 1
);


-- playlist with the highest "order"
SELECT pv1.playlist_id, "order"
FROM z_playlist_video pv1
WHERE "order" ≥ ALL (
    SELECT "order" FROM z_playlist_video
);
```

2.19. SQL S06 L04

- WITH … AS() subquery construction

```sql
SELECT
    SYSDATE - 30 AS back_x_days,
    ADD_MONTHS(SYSDATE, 2) AS back_x_months
FROM dual;


-- channels that uploaded videos last x years
WITH recent_videos AS (
    SELECT * FROM z_video
    WHERE upload_date > ADD_MONTHS(SYSDATE, -2 * 12)
)
SELECT
    channel_id,
    COUNT(*) AS video_count
FROM recent_videos
GROUP BY channel_id;
```

2.20. SQL S07 L01

- INSERT INTO Tab VALUES()
- INSERT INTO Tab (atr, atr) VALUES()
- INSERT INTO Tab AS SELECT …

```sql
-- Simple INSERT
INSERT INTO z_category
VALUES (1, 11, 'Technology');

-- Column-specified INSERT
INSERT INTO z_category(parent_category_id, category_name)
VALUES (11,'Technology');

INSERT INTO z_user (username, first_name, last_name, email)
VALUES ('newuser', 'New', 'User', 'new@user.com');

-- INSERT INTO ... SELECT
INSERT INTO z_playlist (user_id, title, visibility, creation_date)
SELECT user_id, 'Auto Playlist', 'public', SYSDATE
FROM z_user WHERE username = 'newuser';
```

2.21. SQL S07 L02

- UPDATE Tab SET atr = … WHERE condition
- DELETE FROM Tab WHERE atr = …

```sql
-- UPDATE
UPDATE z_video
SET title = 'Updated Title'
WHERE video_id = 1;

-- DELETE
DELETE FROM z_comment
WHERE comment_id = 1;
```

2.22. SQL S07 L03

- DEFAULT, MERGE, Multi-Table Inserts

```sql
-- DEFAULT usage
INSERT INTO z_user (username, first_name, last_name, email, is_deleted)
VALUES ('defaultuser', 'Def', 'User', 'def@user.com', DEFAULT);

-- MERGE example
MERGE INTO z_user u                  -- destination
    USING (
        SELECT
            1 AS user_id,
            'mergeuser' AS username
        FROM DUAL
    ) src                            -- data source
    ON (u.user_id = src.user_id)     -- condition (true → matched, false → not matched)
        WHEN MATCHED THEN
            UPDATE SET u.username = src.username || '_updated'        -- udpate clause
        WHEN NOT MATCHED THEN
            INSERT (user_id, username, first_name, last_name, email)  -- insert clause
            VALUES (src.user_id, src.username, 'M', 'U', 'mu@example.com');



DELETE FROM z_user WHERE username LIKE 'mergeuser%';
SELECT * FROM z_user;
SELECT * FROM z_user WHERE username LIKE 'mergeuser%';
```

2.23. SQL S08 L01

- Objects in databases – Tables, Indexes, Constraint, View, Sequence, Synonym
- CREATE, ALTER, DROP, RENAME, TRUNCATE
- CREATE TABLE (atr DAT TYPE, DEFAULT NOT NULL)
- **ORGANIZATION EXTERNAL, TYPE ORACLE_LOADER, DEFAULT DICTIONARY, ACCESS PARAMETERS, RECORDS DELIMITED BY NEWLINE, FIELDS, LOCATION**

```sql
-- CREATE with constraints
CREATE TABLE z_test_object (
    id NUMBER PRIMARY KEY,
    name NVARCHAR2(100) DEFAULT 'N/A' NOT NULL
);

-- ALTER, RENAME, DROP
ALTER TABLE z_test_object ADD description NVARCHAR2(200);

RENAME z_test_object TO z_test_renamed;

DROP TABLE z_test_renamed;

-- TRUNCATE
INSERT INTO z_test_renamed(id, name)
    SELECT user_id, username
    FROM z_user;

TRUNCATE TABLE z_test_renamed;
```

This cannot be ran on uni's Oracle server. Missing privileges.

```sql
-- create a directory object pointing to the location of the files
CREATE OR REPLACE DIRECTORY ext_tab_dir AS './ext_data';

-- load data from external files
CREATE TABLE x_country_ext (
    country_code      VARCHAR2(5),
    country_name      VARCHAR2(50),
    country_language  VARCHAR2(50)
)
ORGANIZATION EXTERNAL (
    TYPE ORACLE_LOADER
    DEFAULT DIRECTORY ext_tab_dir
    ACCESS PARAMETERS (
        RECORDS DELIMITED BY NEWLINE  -- could be '|' or ';', etc.
        FIELDS TERMINATED BY ','
        MISSING FIELD VALUES ARE NULL
        RECORDS FIXED 512 FIELDS      -- 512 bytes per record
        (
            country_code      CHAR(5),
            country_name      CHAR(50),
            country_language  CHAR(50)
        )
    )
    LOCATION (
        'Countries1.txt',
        'Countries2.txt'
    )
);

-- Countries1.txt
/*
ENG,England,English
SCO,Scotland,English
IRE,Ireland,English
WAL,Wales,Welsh
*/

-- Countries2.txt
/*
FRA,France,French
GER,Germany,German
USA,Unites States of America,English
*/
```

2.24. SQL S08 L02

- TIMESTAMP, TIMESTAMP WITH TIME ZONE, TIMESTAMP WITH LOCAL TIMEZONE
- INTERVAL YEAR TO MONTH, INTERVAL DAY TO SECOND
- CHAR, VARCHAR2, CLOB
- about NUMBER
- about BLOB

```sql
-- Advanced data types
CREATE TABLE z_types_demo (
    n       NUMBER,
    c       CHAR(10),
    vc      VARCHAR2(100),
    ts      TIMESTAMP,
    ts_tz   TIMESTAMP WITH TIME ZONE,
    ts_ltz  TIMESTAMP WITH LOCAL TIME ZONE,
    iv1     INTERVAL YEAR TO MONTH,   -- stores intervals using year and month
    iv2     INTERVAL DAY TO SECOND,   -- stores intervals using days, hours, minutes,
                                      -- and seconds including fractional seconds
    txt     CLOB, -- character large object
    bin     BLOB -- binary large object
);

SELECT * FROM z_types_demo;

DROP TABLE z_types_demo;

CREATE TABLE z_demo_student (
    years_of_undergrad INTERVAL YEAR TO MONTH
);

INSERT INTO z_demo_student VALUES (INTERVAL '10-2' YEAR TO MONTH);
INSERT INTO z_demo_student VALUES (INTERVAL '2-3' YEAR TO MONTH);
INSERT INTO z_demo_student VALUES (INTERVAL '9' MONTH);

SELECT * FROM z_demo_student;

DROP TABLE z_demo_student;
```

2.25. SQL S08 L03

- ALTER TABLE (ADD, MODIFY, DROP), DROP, RENAME
- FLASHBACK TABLE Tab TO BEFORE DROP (view USER_RECYCLEBIN)
- DELETE, TRUNCATE
- COMMENT ON TABLE
- SET UNUSED

```
-- ALTER operations
ALTER TABLE z_demo_student ADD email NVARCHAR2(100);
ALTER TABLE z_demo_student MODIFY email NVARCHAR2(300);
ALTER TABLE z_demo_student DROP COLUMN email;

CREATE TABLE z_demo_student (years_of_undergrad INTERVAL YEAR TO MONTH);

INSERT INTO z_demo_student VALUES (INTERVAL '10-2' YEAR TO MONTH);
INSERT INTO z_demo_student VALUES (INTERVAL '2-3' YEAR TO MONTH);
INSERT INTO z_demo_student VALUES (INTERVAL '9' MONTH);

DROP TABLE z_demo_student;

-- FLASHBACK
FLASHBACK TABLE z_demo_student TO BEFORE DROP;

-- COMMENT
COMMENT ON TABLE z_demo_student IS 'Holds student info and stuff';

-- SET UNUSED
-- source: https://forums.oracle.com/ords/apexds/post/set-unused-and-drop-column-4001
-- The SET UNUSED option marks one or more columns as unused so
-- that they can be dropped when the demand on system resources
-- is lower. Specifying this clause does not actually remove the
-- target columns from each row in the table (that is, it does not
-- restore the disk space used by these columns).Therefore, the
-- response time is faster than if you executed the DROP clause.
-- Unused columns are treated as if they were dropped, even though
-- their column data remains in the tables rows. After a column
-- has been marked as unused, you have no access to that column.
ALTER TABLE z_demo_student SET UNUSED (email);
```

2.26. SQL S10 L01

- CREATE TABLE (NOT NULL AND UNIQUE constraint)
- CREATE TABLE Tab AS SELECT …
- Own vs. system naming CONSTRAINT conditions

```sql
-- Table with named constraints
CREATE TABLE z_example_constraints (
    id NUMBER CONSTRAINT pk_example PRIMARY KEY,
    name NVARCHAR2(100) CONSTRAINT nn_example_name NOT NULL,
    code NVARCHAR2(10) CONSTRAINT uq_example_code UNIQUE
);

-- CREATE TABLE AS SELECT
CREATE TABLE z_user_copy AS
    SELECT *
    FROM z_user;

SELECT *
FROM z_user_copy;

DROP TABLE z_user_copy;
```

2.27. SQL S10 L02

- CONSTRAINT – NOT NULL, UNIQUE, PRIMARY KEY, FOREIGN KEY (atr REFERENCES Tab(atr)), CHECK
- Foreign keys, ON DELETE, ON UPDATE, RESTRICT, CASCADE, etc.

```sql
-- Foreign key with ON DELETE CASCADE
CREATE TABLE z_fk_demo (
    id NUMBER PRIMARY KEY,
    user_id NUMBER,
    CONSTRAINT fk_demo_user
        FOREIGN KEY (user_id) REFERENCES z_user(user_id)
        ON DELETE CASCADE
);

-- CHECK constraint
ALTER TABLE z_fk_demo ADD CONSTRAINT ck_demo_id CHECK (id > 0);
```

2.28. SQL S10 L03

- about USER_CONSTRAINTS

```sql
-- Inspect user constraints of table "z_user"
SELECT
    constraint_name,
    constraint_type,
    table_name
FROM user_constraints
WHERE table_name = 'Z_USER';
```

2.29. SQL S11 L01

- CREATE VIEW
  - about FORCE, NOFORCE
  - WITCH CHECK OPTION
  - WITH READ ONLY
- about Simple vs. Compex VIEW

```sql
-- Simple view
CREATE OR REPLACE VIEW v_active_users AS
    SELECT user_id, username, email
    FROM z_user
    WHERE is_deleted = 0;

-- View with CHECK OPTION
--    any INSERT or UPDATE done through the view must result
--    in rows that are still visible in the view
--    that is, they must satisfy the view's WHERE condition.
CREATE OR REPLACE VIEW v_public_playlists AS
    SELECT *
    FROM z_playlist
    WHERE visibility = 'public'
WITH CHECK OPTION;

-- Read-only view
CREATE OR REPLACE VIEW v_readonly AS
    SELECT video_id, title
    FROM z_video
WITH READ ONLY;
```

2.30. SQL S11 L03

- INLINE VIEW Subquery in the form of a table SELECT atr FROM (SELECT * FROM Tab) alt_tab

```sql
-- Inline view
SELECT alt.title
FROM (
    SELECT title
    FROM z_video
    WHERE is_deleted = 0
) alt;
```

2.31. SQL S12 L01

- CREATE SEQUENCE name INCREMENT BY n START WITH m, (NO)MAXVALUE, (NO)MINVALUE, (NO)CYCLE, (NO)CACHE
- about ALTER SEQUENCE

```sql
-- Sequence creation and alteration
CREATE SEQUENCE seq_demo_student_id
    INCREMENT BY 1
    START WITH 100
    NOMAXVALUE
    NOCYCLE
    NOCACHE;

ALTER SEQUENCE seq_demo_student_id INCREMENT BY 5;

CREATE TABLE z_demo_student (
    id NUMBER DEFAULT seq_demo_student_id.nextval,
    name VARCHAR2(40 CHAR)
);

INSERT INTO z_demo_student(name) VALUES ('john');
INSERT INTO z_demo_student(name) VALUES ('doe');

SELECT * FROM z_demo_student;

DROP SEQUENCE seq_demo_student_id;

DROP TABLE z_demo_student;
```

2.32. SQL S12 L02

- CREATE INDEX, PRIMARY KEY, UNIQUE KEY, FOREIGN KEY

```sql
-- Index creation
CREATE INDEX idx_z_demo_student_name ON z_demo_student(name);
```

2.33. SQL S13 L01

- GRANT … ON … TO … PUBLIC
- about REVOKE
- What rights can be assigned to which objects? (ALTER, DELETE, EXECUTE, INDEX,

INSERT, REFERENCES, SELECT, UPDATE) – (TABLE, VIEW, SEQUENCE, PROCEDURE)

```sql
-- Granting privileges
GRANT SELECT, INSERT, UPDATE ON z_user TO PUBLIC;

GRANT DELETE ON z_video TO some_user;

-- Revoking privileges
REVOKE SELECT, INSERT, UPDATE ON z_user FROM PUBLIC;
```

2.34. SQL S13 L03

- Regular expressions
- REGEXP_LIKE, REGEXP_REPLACE, REGEXP_INSTR, REGEXP_SUBSTR, REGEXP_COUNT

```sql
-- Regular expressions
SELECT email
FROM z_user
WHERE REGEXP_LIKE(email, '^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}$');

SELECT REGEXP_REPLACE(
    '123-456-7890',
    '[^0-9]',
    'non_number'
) FROM DUAL;


-- returns index 3 because number '1' is at that position
SELECT REGEXP_INSTR(
    'aa1b2c3',
    '[0-9]'
) FROM DUAL;


SELECT REGEXP_SUBSTR(
    'abc123def456ghi789jkl012',
    '[0-9]+',
    10,                        -- starting pos in source string
    3                          -- occurrence
) FROM DUAL;

-- 8
SELECT REGEXP_COUNT(
    'a1b2c3d45e67f8g',
    '[0-9]'
) FROM DUAL;

-- 6
SELECT REGEXP_COUNT(
    'a1b2c3d45e67f8g',
    '[0-9]+'
) FROM DUAL;
```

2.35. SQL S14 L01

- Transactions, COMMIT, ROLLBACK, SAVEPOINT

```sql
SELECT * FROM z_user WHERE user_id = 1;

-- Transaction control
DECLARE
    v_success INT := 0;
BEGIN
    SAVEPOINT before_update;

    -- do some transaction stuff
    UPDATE z_user
    SET email = 'new@email.com'
    WHERE user_id = 1;

    -- throw error if something goes wrong
    IF v_success = 0 THEN
        RAISE_APPLICATION_ERROR(50001, 'Transaction failed :(');
    END IF;

    -- if nothing went wrong, commit the changes
    COMMIT;
    DBMS_OUTPUT.PUT_LINE('Transaction successful – changes commited');

EXCEPTION
    WHEN OTHERS THEN
        -- if anything goes wrong, rollback
        ROLLBACK TO before_update;
        DBMS_OUTPUT.PUT_LINE('Transaction failed – rolling back changes');
END;
```

2.36. SQL S15 L01

- Alternative join notation without JOIN with join condition in WHERE
- Left and right connection using atrA = atrB (+)

```sql
-- Join using WHERE (legacy style)

SELECT *
FROM
    z_user u,
    z_channel c
WHERE u.user_id = c.user_id;


-- LEFT and RIGHT JOIN using (+) operator

SELECT *
FROM z_video v
LEFT JOIN z_video_view vv ON v.video_id = vv.video_id
WHERE vv.VIDEO_VIEW_ID IS NULL;
-- equivalent to
SELECT *
FROM
    z_video v,
    z_video_view vv
WHERE v.video_id = vv.video_id(+) -- the side with (+) is 'added'
  AND vv.VIDEO_VIEW_ID IS NULL;



SELECT *
FROM z_video_view vv
RIGHT JOIN z_video v ON vv.video_id = v.video_id;
-- equivalent to
SELECT *
FROM
    z_video_view vv,
    z_video v
WHERE v.video_id = vv.video_id(+);
```